

- Assignment: 8.2 Time Series Modeling
  - Name: Barath Anandaraman
  - Course: DSC630-T301
  - Week8: Time Series
  - Date: 04/30/2025
- 

## Load necessary packages

```
In [37]: # Load libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from sklearn.metrics import mean_squared_error
```

## Load the US Retail Sales dataset

```
In [2]: retail_df = pd.read_csv('us_retail_sales.csv')
retail_df.head()
```

```
Out[2]:
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC
0	1992	146925	147223	146805	148032	149010	149800	150761.0	151067.0	152588.0	153521.0	153583.0	155614.0
1	1993	157555	156266	154752	158979	160605	160127	162816.0	162506.0	163258.0	164685.0	166594.0	168161.0
2	1994	167518	169649	172766	173106	172329	174241	174781.0	177295.0	178787.0	180561.0	180703.0	181524.0
3	1995	182413	179488	181013	181686	183536	186081	185431.0	186806.0	187366.0	186565.0	189055.0	190774.0
4	1996	189135	192266	194029	194744	196205	196136	196187.0	196218.0	198859.0	200509.0	200174.0	201284.0

The original dataset has years in rows and months as columns. We'll melt it to long format for time series analysis

```
In [3]: # Convert wide to long format
retail_melted = retail_df.melt(id_vars=['YEAR'], var_name='Month', value_name='Sales')
# Map month names to month names
```

```

month_num = {
    'JAN': 1, 'FEB': 2, 'MAR': 3, 'APR': 4, 'MAY': 5, 'JUN': 6,
    'JUL': 7, 'AUG': 8, 'SEP': 9, 'OCT': 10, 'NOV': 11, 'DEC': 12
}
retail_melted['Month_Num'] = retail_melted['Month'].map(month_num)
# Create a proper datetime column
retail_melted['Date'] = pd.to_datetime(dict(year=retail_melted['YEAR'], month=retail_melted['Month_Num'], day=1))

# Sort the table by date
retail_melted = retail_melted.sort_values('Date')

```

In [4]: `retail_melted.head()`

Out[4]:

	YEAR	Month	Sales	Month_Num	Date
0	1992	JAN	146925.0	1	1992-01-01
30	1992	FEB	147223.0	2	1992-02-01
60	1992	MAR	146805.0	3	1992-03-01
90	1992	APR	148032.0	4	1992-04-01
120	1992	MAY	149010.0	5	1992-05-01

Check for NAs and drop them

In [27]: `# Check for NAs`  
`print(retail_melted.columns[retail_melted.isna().any()].tolist())`  
`['Sales']`

In [28]: `retail_melted['Sales'].isna().sum()`

Out[28]: `np.int64(6)`


Column Sales has 6 rows of NA and has to be dropped

In [30]: `retail_melted = retail_melted.dropna(subset=['Sales'])`

---

Step1. Plot the data with proper labeling and make some observations on the graph.

```
In [32]: plt.figure(figsize=(10,6))
plt.plot(retail_melted['Date'], retail_melted['Sales'])
plt.title('Monthly US Retail Sales (1992 - 2021)')
plt.xlabel('Date')
plt.ylabel('Retail Sales (in Millions)')
plt.grid(True)
plt.show()
```



Achieved so far:

Loaded the data in dataframe

Reshaped the data from wide format to long format

Dropped NA rows

Plotted monthly retail sales from 1992 to 2021

### Observations from graph

Sales show a dip around 2008 and 2020

- 2008, could be due to the housing crisis
- 2020, could be due to COVID crisis

There is a steady upward trend over the years - retail sales generally increase over time

---

Step2. Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set.

```
In [33]: # Per requirement Define Test set: July 2020 till June 2021
test = retail_melted[(retail_melted['Date'] >= '2020-07-01') & (retail_melted['Date'] <= '2021-06-01')]
# Define Train set as everything else
train = retail_melted[~retail_melted['Date'].isin(test['Date'])]
```

---

Step3. Use the training set to build a predictive model for the monthly retail sales.

```
In [34]: # Fit model using training data
model = ExponentialSmoothing(train['Sales'],
                             trend='add',
                             seasonal='add',
                             seasonal_periods=12).fit()
```

```
/opt/anaconda3/lib/python3.12/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: An unsupported index was provided. As a result, forecasts cannot be generated. To use the model for forecasting, use one of the supported classes of index.
```

```
self._init_dates(dates, freq)
/opt/anaconda3/lib/python3.12/site-packages/statsmodels/tsa/holtwinters/model.py:918: ConvergenceWarning: Optimization failed to converge. Check mle_retvals.
warnings.warn(
```


---

Step4. Use the model to predict the monthly retail sales on the last year of data.

```
In [35]: # Predict on test period
predictions = model.forecast(steps=12)

# Plot actual vs predicted
plt.figure(figsize=(10,6))
plt.plot(train['Date'], train['Sales'], label = 'Training Data')
plt.plot(test['Date'], test['Sales'], label = 'Actual Sales (Test Set)', color = 'green')
plt.plot(test['Date'], predictions, label = 'Predicted Sales', color = 'red' , linestyle='--')
plt.title('Retail Sales Prediction')
plt.xlabel('Date')
plt.ylabel('Retail Sales (in Millions)')
plt.legend()
plt.grid(True)
plt.show()
```

```
/opt/anaconda3/lib/python3.12/site-packages/statsmodels/tsa/base/tsa_model.py:837: ValueWarning: No supported index is
available. Prediction results will be given with an integer index beginning at `start`.
    return get_prediction_index(
/opt/anaconda3/lib/python3.12/site-packages/statsmodels/tsa/base/tsa_model.py:837: FutureWarning: No supported index is
available. In the next version, calling this method in a model without a supported index will result in an exception.
    return get_prediction_index(
```



---

Step5. Report the RMSE of the model predictions on the test set.

```
In [38]: # Calculate RMSE
rmse = np.sqrt(mean_squared_error(test['Sales'], predictions))
rmse
```

```
Out[38]: np.float64(45156.97171858635)
```

---

Achieved so far:

Model used: Holt-Winters Exponential Smoothing (captures both trend and seasonality)

Prediction Period: July 2020 - June 2021

Root Mean Squared Error on Test Set: 45,156.97, this means the typical prediction error is around \$45,157 million

## Final Summary

Training data: January 1992 to June 2020

Test data: July 2020 to June 2021


Observation: The model closely tracks the trend, but some deviation occurs possibly due to COVID-era volatility

---

## Additionally check for Seasonal Volatility using Box Plot

```
In [40]: import seaborn as sns
# Ensure month order
df_melted = retail_melted.copy()
df_melted['Month'] = pd.Categorical(df_melted['Month'],
                                   categories = ['JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN',
                                                  'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC'],
                                   ordered = True)

plt.figure(figsize=(14,6))
sns.boxplot(x='Month', y='Sales', data = df_melted)
plt.title('Seasonal Volatility: Monthly Retail Sales Spread by Month')
plt.ylabel('Sales (in Millions)')
plt.grid(True)
plt.show()
```



Noticed certain months March, May, June have higher retail sales this could be seasonality

In [ ]: